

Distributed Matrix Computation (DMaC)

A Distributed System for Matrix Multiplication

Ryan Bazzell & Jiwoon Yim

March 17, 2025

Introduction

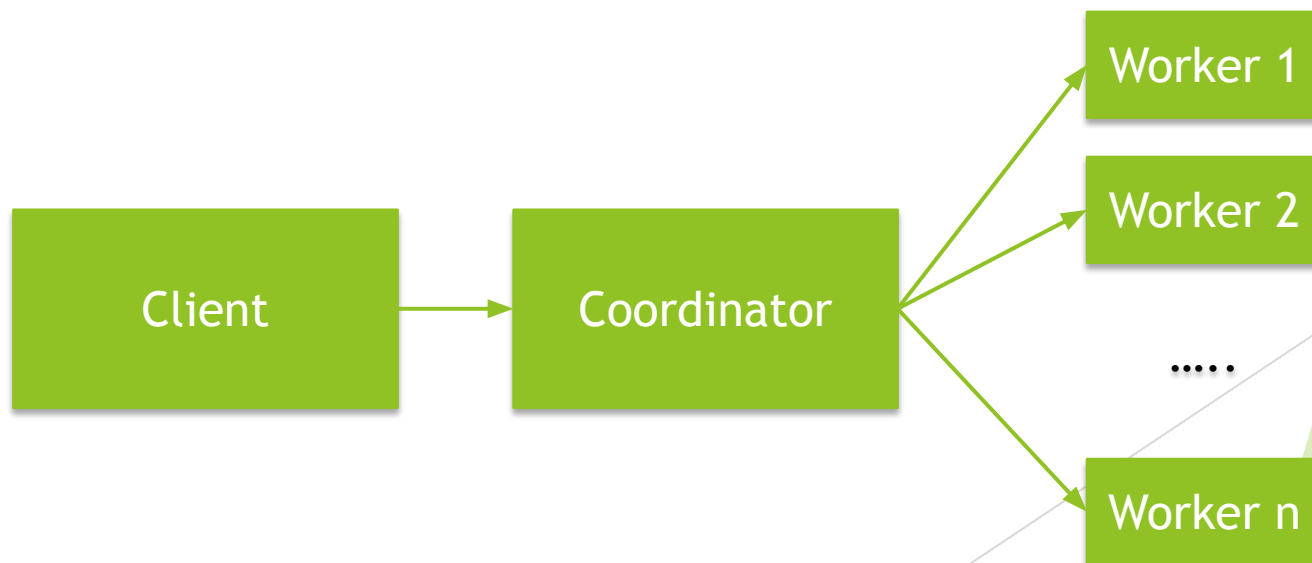
- ▶ Distributed matrix multiplication
 - ▶ Strassen's Algorithm
- ▶ Heterogeneous nodes
- ▶ Coordinator-worker architecture
- ▶ Client for User-Interface
- ▶ Scalable and transparent

Motivation

- ▶ Single-threaded shortcomings
- ▶ Time consuming large-scale matrix multiplications
- ▶ Improve performance
- ▶ Provide scalability

System Architecture

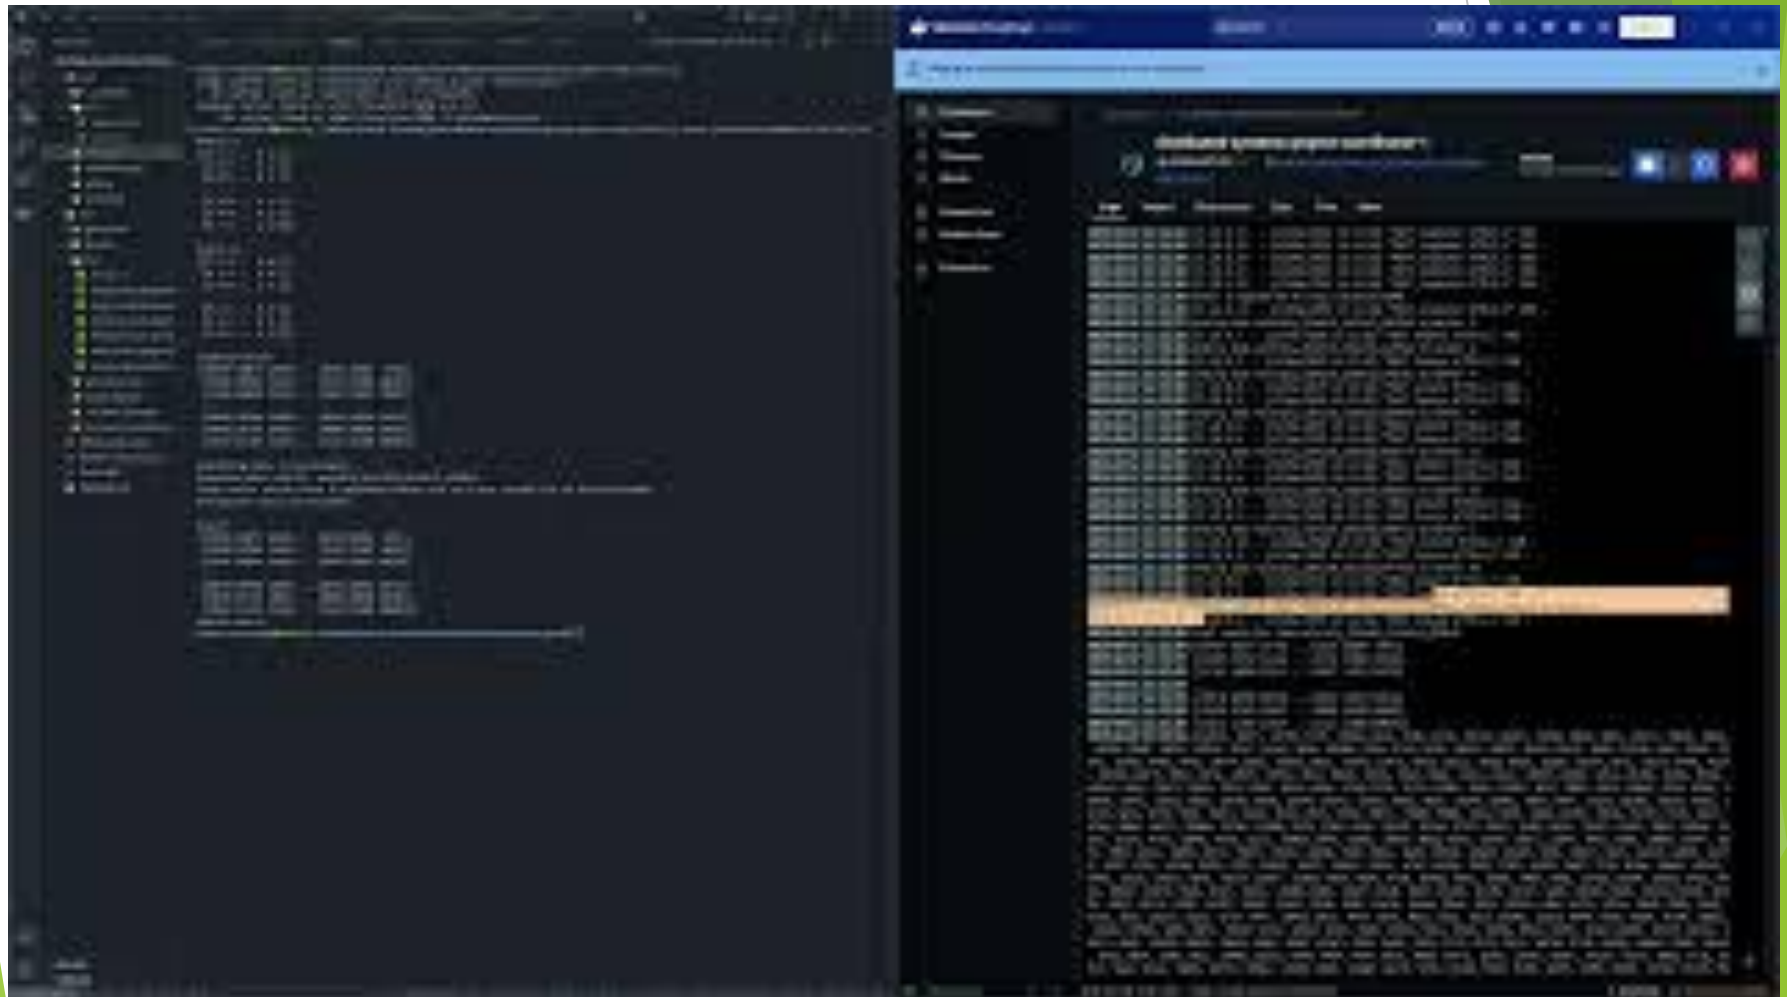
- ▶ Client: Submits jobs and verifies results
- ▶ Coordinator: Distributes tasks, collects results
- ▶ Workers: Compute using Strassen's Algorithm



Implementation Details

- ▶ Failure handling with retry mechanisms
- ▶ Python with Flask, NumPy, and threading
- ▶ Docker simulates a distributed environment
- ▶ **Strassen's Algorithm** for matrix multiplication ($O(n^{2.807})$ vs. $O(n^3)$)

Demo




Simulation & Testing


- ▶ Performance benchmarks and test cases
 - ▶ Simulated heterogeneous setup with Docker (Alpine, Debian, Fedora)
 - ▶ Three scales: Small (4 workers), Medium (10 workers), Large (20 workers)
 - ▶ Compared local vs. distributed performance
- ▶ Metrics
 - ▶ Accuracy
 - ▶ Performance
 - ▶ Scalability
 - ▶ Transparency

Results - Diagnostics


Small Scale Diagnostic



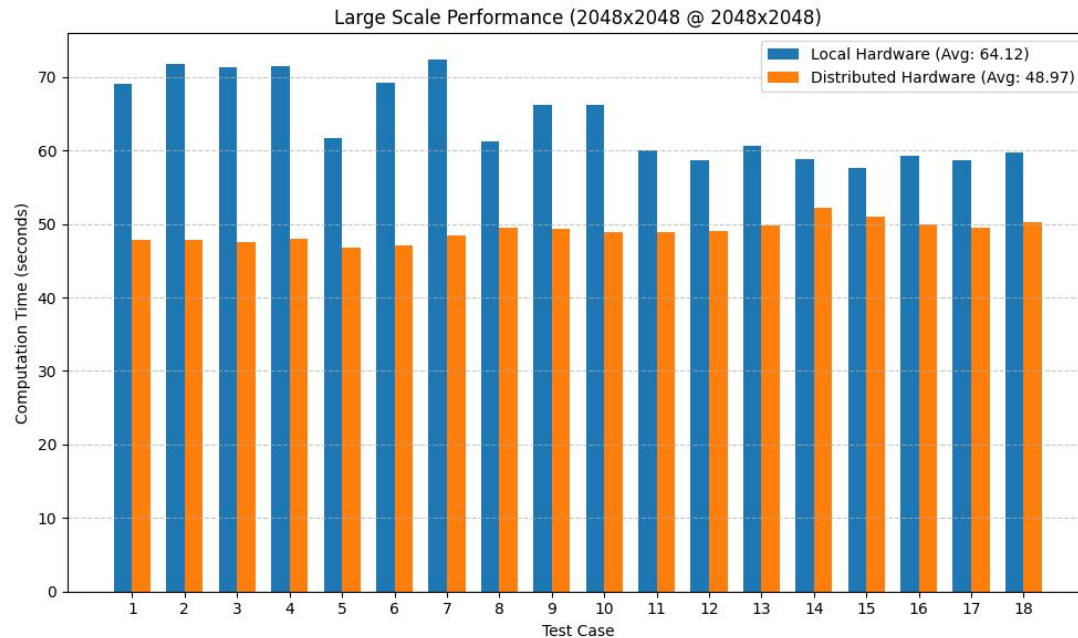
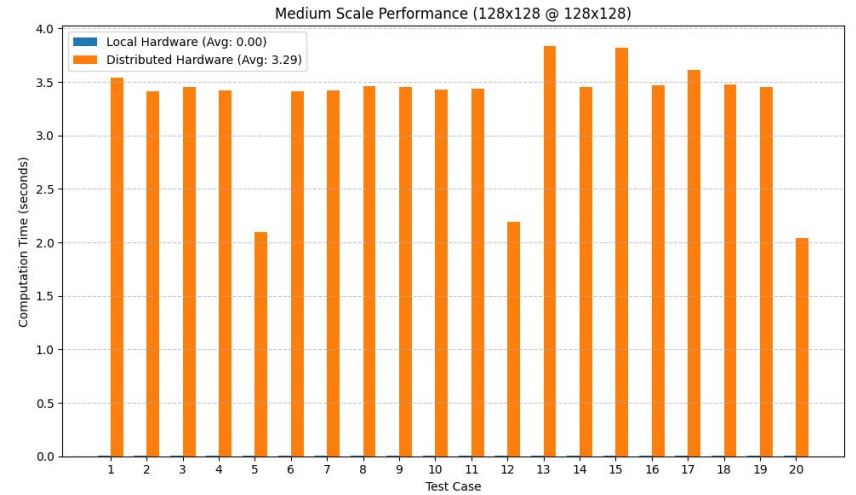

Medium Scale Diagnostic



Large Scale Diagnostic



Results - Performance



Distributed System Principles

- ▶ Scalability: Works across different worker counts
- ▶ Transparency: Hides distribution from user
- ▶ Resource Sharing: Tasks split across nodes



Lessons Learned

- ▶ Communication Overhead
- ▶ Single Point of Failure
- ▶ Effectiveness of Programming Tools
- ▶ Efficient Distributed Algorithms
- ▶ Fault Tolerance

Future Work

- ▶ Addressing the Single Point of Failure
- ▶ Improving Communication Efficiency
- ▶ Implement Worker Fault Tolerance

Conclusion

- ▶ Distributed computing for matrix multiplications
- ▶ Offers scalability, though overhead remains a challenge
- ▶ Future improvements can enhance real-world applicability

References

- ▶ API – Flask Documentation (3.0.x). 2024. url: <https://flask.palletsprojects.com/en/stable/api/>.
- ▶ Ankur Mallick, Malhar Chaudhari, and Gauri Joshi. “Fast and Efficient Distributed Matrix-vector Multiplication Using Rateless Fountain Codes”. In: ICASSP 2019 - 2019 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP). 2019, pp. 8192-8196. doi: 10.1109/ICASSP.2019.8682347.
- ▶ Oded Schwartz and Noa Vaknin. “Pebbling Game and Alternative Basis for High Performance Matrix Multiplication”. In: SIAM Journal on Scientific Computing 45.6 (2023), pp. C277-C303. doi: 10.1137/22M1502719. eprint: <https://doi.org/10.1137/22M1502719>. url: <https://doi.org/10.1137/22M1502719>.
- ▶ Volker Strassen. “Gaussian elimination is not optimal”. In: Numerische Mathematik 13.4 (Aug. 1969), pp. 354-356. doi:10.1007/bf02165411. url: <https://doi.org/10.1007/bf02165411>.
- ▶ Serhii Zybin, Vladimir Khoroshko, Volodymyr Maksymovych, and Ivan Opirskyy. “Effective Distribution of Tasks in Multiprocessor and Multi-Computers Distributed Homogeneous Systems”. In: International Journal of Computing (June 2021), pp. 211-220. doi: 10.47839/ijc.20.2.2168.